

9 - JDBC

[< Previous](#) | [Next >](#)

§9.1 Fundamentals

- A JDBC driver implementation must provide implementation classes for the following interfaces: `java.sql.Driver`, `java.sql.Connection`, `java.sql.Statement` and `java.sql.ResultSet`.
- Provided by the JDBC API are: `java.sql.DriverManager`, `java.sql.SQLException` and `java.sql.Date`.
- Objects of the following classes/interfaces are independent of the implementation of a JDBC driver: `javax.sql.RowSet`, `java.sql.Date`, `java.sql.Time` and `java.sql.SQLException`.
- Applications no longer need to explicitly load JDBC drivers using `Class.forName()`. JDBC 4.0 Drivers must include the file `META-INF/services/java.sql.Driver`. This file contains the name of the JDBC drivers implementation of `java.sql.Driver`. Therefore, no java code is necessary to load the driver classes.
- When the method `getConnection` is called `DriverManager.getConnection()`, the `DriverManager` will attempt to locate a suitable driver from amongst those loaded at initialization and those loaded explicitly using the same classloader as the current applet or application.
- The method signatures of the `DriverManager.getConnection()` method:
 - `public static Connection getConnection(String url)`
 - `public static Connection getConnection(String url, String user, String password)`
 - `public static Connection getConnection(String url, java.util.Properties info)`
- You can set the user and password in a `Properties` object via the `setProperty()` method. The correct property names are "user" and "password". `jdbc.driver` is also an invalid property name but it is not an error to use it.
- `DriverManager.registerDriver` is a valid method but it takes `java.sql.Driver` instance and not a `String`. This method is used by the Driver class to register itself with the `DriverManager`. It need not be called by the application programmer. In JDBC 4.0, if you have the jar file that implements the Driver in the classpath, the Driver is automatically registered.
- There is no `setClientInfo` or `connect` method in `Connection`.
- When a connection is created, it is in auto-commit mode. i.e. auto-commit is enabled. This means that each individual SQL statement is treated as a transaction and is automatically committed right after it is completed.
- When auto-commit has been disabled in the given code (by calling `c.setAutoCommit(false)`), you have to explicitly commit the transaction to commit the changes to the database. The regular way to do this is to **call `con.commit()`**. Notice that commit method does not take any arguments.
- **There is no commit method in Statement.**
- Another way is to utilize the side effect of changing the auto-commit mode of the connection. If the `setAutoCommit` method is called during a transaction and the auto-commit mode is changed, the transaction is committed. If `setAutoCommit` is called and the auto-commit mode is not changed, the call is a no-op.
- The way to allow two or more statements to be grouped into a transaction is to disable the auto-commit mode. Since it is enabled by default, you have to explicitly disable it after creating a connection by calling `con.setAutoCommit(false)`;
- Calling `connection.rollback()` without passing in a savepoint, will cause the whole transaction to rollback.
- If the `connection.close()` method is called and there is an active transaction, the results are implementation-defined. Therefore, calling this method does not guarantee that the active transaction will be committed.
- As per Section 6.2 of JDBC 4.1 Specification: A JDBC API implementation must support Entry Level SQL92 plus the SQL command Drop Table. Entry Level SQL92 represents a "floor" for the level of SQL that a JDBC API implementation must support. Access to features based on SQL99 or SQL:2003 should be provided in a way that is compatible with the relevant part of the SQL99 or SQL:2003 specification.

- A `java.sql.Connection` object implements a session with the database and provides methods to commit and rollback transactions.
- A `java.sql.Statement` object provides various execute methods to execute queries on the database.
- All queries run on the database.
- All queries executed on the database have to run in native SQL, whether fired through `Statement`, `PreparedStatement`, or from within a stored procedure.

ResultSet:

- A `java.sql.ResultSet` object provides a cursor to fetch data from the database.
- The numbering of columns in a `ResultSet` starts with 1. The value of a field with type `INT`, can still be retrieved using `getString()`. Note that if a field is of type `VARCHAR` and if you try to retrieve the value using say `getInt()` or `getDouble()`, it may throw an exception at runtime if the value cannot be parsed into an `Integer` or `Double`.
- `getString` retrieves the value of the designated column in the current row of this `ResultSet` object as a `String`. If the value is `SQL NULL`, the value returned is `null`.
- `rs.next()` returns `true` only if there is a row left to be processed in the `ResultSet`. It moves the cursor in front of the next row and returns `true`. If there is no row left, it returns `false`. Thus, when you use `rs.next()` in a while loop, you are basically iterating through the rows returned by the query one by one.
- `Statement` has a `setMaxRows` method that limits the total number of rows returned by the `ResultSet`. If the query returns more rows than the limit, the extra rows are silently (i.e. without any exception) are ignored.
- `java.sql.ResultSetMetaData` gives you the information about the result of executing a query. You can retrieve this object by calling `getMetaData()` on `ResultSet`. `ResultSetMetaData` contains several methods that tell you about the `ResultSet`. Some important methods are: `getColumnCount()`, `getColumnName(int col)`, `getColumnLabel(int col)`, and `getColumnType(int col)`. Remember that the column index starts from 1.
- If a `Statement` object from which a `ResultSet` was retrieved is closed, its current `ResultSet` object, if one exists, is also closed. If the `ResultSet` object that you are iterating is closed, you will get an exception saying:
`Exception in thread "main" java.sql.SQLException: ResultSet not open`.
- JDBC 2.0 allows you to use `ResultSet` object to update an existing row and even insert new row in the database. For both the cases, the `ResultSet` must be updatable, which can be achieved by passing `ResultSet.CONCUR_UPDATABLE` while creating a `Statement` object:
`stmt = con.createStatement(ResultSet.TYPE_SCROLL_INSENSITIVE, ResultSet.CONCUR_UPDATABLE);`
or
`stmt = con.createStatement(ResultSet.TYPE_SCROLL_SENSITIVE, ResultSet.CONCUR_UPDATABLE);`
- To update an existing row:
 1. First, go to the row you want to update. You can either iterate through a `ResultSet` to reach a particular row or just call `rs.absolute(int rowNumber)`.
 2. Now update the columns of the `ResultSet` with required values using `rs.updateXXX(columnNumber, value)` or `rs.updateXXX(columnName, value)` methods.
 3. Call `rs.updateRow();` If you call `rs.refreshRow()` without calling `updateRow()`, the updates will be lost.
- To insert a new Row:
 1. Call `rs.moveToInsertRow();` first. You can't insert a row without calling this method first.
 2. Use `rs.updateXXX` methods to update all column values. You must set values for all the columns.
 3. Call `rs.insertRow();`
 4. Call `rs.moveToCurrentRow();` to go back to the row where you were before calling `moveToInsertRow`.
- **IMPORTANT:** The exam will test you on implications of calling various methods out of sequence. For example, what

happens when you call `insertRow` without first calling `moveToInsertRow`? (An `SQLException` will be thrown.) or what happens when you call `refreshRow` without first calling `updateRow`? (No exception but updates will be lost.).

CallableStatement vs. PreparedStatement:

PreparedStatement

- A `PreparedStatement` is used for SQL statements that are executed multiple times with different values. For example, if you want to insert several values into a table, one after another, it is a lot easier with `PreparedStatement`:

```
ps = c.prepareStatement("INSERT INTO STUDENT VALUES (?, ?)");
//This is created only once
//Once created, the PreparedStatement is compiled automatically.
ps.setInt(1, 111);
ps.setString(2, "Bob");
ps.executeUpdate();

//Now change the parameter values and execute again.
ps.setInt(1, 112);
ps.setString(2, "Cathy");
ps.executeUpdate();
```

- `PreparedStatement` offers better performance when the same query is to be run multiple times with different parameter values.
- `PreparedStatement` has specific methods for additional SQL column type such as `setBlob(int parameterIndex, Blob x)` and `setClob(int parameterIndex, Clob x)`.

CallableStatement

- A `CallableStatement` is meant for executing a stored procedure, which has already been created in the database. For example:

```
//computeMatrixForSales is a stored procedure that has already been created in the database.
callableStatement = connection.prepareCall("{call computeMatrixForSales(?)}");
callableStatement.setInt(1, 1000);
callableStatement.executeUpdate();
```

- One advantage of `CallableStatement` is that it allows IN/OUT parameters.
- A `CallableStatement` is the only way for a JDBC program to execute stored procedures in the database if the procedure has in and out parameters.